

IN THE NAME OF ALLAH

Linear Regression

Implementing a linear regression for a real data set
(body measurements) with scikit-learn (sklearn) library in python

Seyed Mahdi Mahdavi Mortazavi - 40030490

Numerical Computing Methods

July 2023

By theMHD

Introduction

What is the linear regression?

Linear regression analysis is used to predict value of a variable (dependent variable) based on the value of another variables (independent variables) in a data set.

Linear regression model takes this form:

$$y_i = a_0 + a_1x_1 + a_2x_2 + a_3x_3 + \dots + a_nx_n + \varepsilon_i = a_0 + \sum_{i=[1,n]} a_ix_i + \varepsilon_i$$

$y \rightarrow$ dependent variable / $x \rightarrow$ independent variable / a_1 to $a_n \rightarrow$ coefficients of independent variables / $a_0 \rightarrow$ intercept of linear equation / $\varepsilon_i \rightarrow$ error of prediction

The goal of this project is to predict a person' age using his/her body measurements (with a real data set) ...

Let's go code!

Step 1

Get some general information from our data set file (part1)

.head() method and it's results (it shows the first 5 rows of data set file as it's header)

5 rows and 13 columns

```
print("\n||| ----- Step 1: General file info ----- |||")
print('1) The first 5 rows -----')
print(body_measures.head())
print('\n2) File information -----')
print(body_measures.info())
print('\n3) Statistical information -----')
print(body_measures.describe())
```

```
||| ----- Step 1: General file information ----- |||
```

```
1) The first 5 rows -----
```

	Gender	Age	HeadCircumference	...	WaistToKnee	LegLength	TotalHeight
0	1	30	22	...	25	22	52
1	1	28	19	...	25	20	56
2	2	27	21	...	14	18	53
3	1	29	20	...	20	21	45
4	2	28	16	...	32	13	47

```
[5 rows x 13 columns]
```

```

2) File information -----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 716 entries, 0 to 715
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                 716 non-null   int64
1   Age                    716 non-null   int64
2   HeadCircumference      716 non-null   int64
3   ShoulderWidth          716 non-null   int64
4   ChestWidth             716 non-null   int64
5   Belly                  716 non-null   int64
6   Waist                  716 non-null   int64
7   Hips                   716 non-null   int64
8   ArmLength              716 non-null   int64
9   ShoulderToWaist        716 non-null   int64
10  WaistToKnee            716 non-null   int64
11  LegLength              716 non-null   int64
12  TotalHeight            716 non-null   int64
dtypes: int64(13)
memory usage: 72.8 KB
None

```

Step 2 (.info())

Get some general information from our data set file (part2)

More with code! ...

```

3) Statistical information -----

```

	Gender	Age	...	LegLength	TotalHeight
count	716.000000	716.000000	...	716.000000	716.000000
mean	1.452514	15.340782	...	26.833799	48.118715
std	0.498088	11.831501	...	7.925988	12.156722
min	1.000000	1.000000	...	9.000000	19.000000
25%	1.000000	7.000000	...	21.000000	40.000000
50%	1.000000	11.000000	...	26.000000	48.000000
75%	2.000000	21.000000	...	32.000000	55.000000
max	2.000000	68.000000	...	50.000000	89.000000

```

[8 rows x 13 columns]

```

Step 3 (.describe())

Get some statistical information about our data set file ...
Such as mean, std (root of variance), min, max and ...


```

55 print("\n||| ----- Step 2: Make regression line ----- |||")
56 x_list = body_measures[
57     ["Gender", "HeadCircumference", "ShoulderWidth", "ChestWidth", "Belly", "Waist", "Hips",
58     "ArmLength", "ShoulderToWaist", "WaistToKnee", "LegLength", "TotalHeight"]]
59 y_list = body_measures["Age"]
60 reg_res = LinearRegression()
61 reg_res.fit(x_list, y_list)
62
63 i = 0 # index of coefficients
64 print("The coefficient of variables are shown in bottom...")
65 for header in body_measures.columns:
66     if header != "Age":
67         print(f"{header}: {reg_res.coef_[i]}")
68         i += 1
69
70 print("-----")
71 print("Intercept of regression line is: ", reg_res.intercept_)
72
73 # Predict Ages and calculate error -----
74 predicted_y = reg_res.predict(x_list) # predicated y list (y is list Ages)
75 rmse = np.sqrt(np.mean(np.square(predicted_y - y_list))) # root mean square error
76 print("Average prediction error: ", rmse, " ~= ", round(rmse))

```

With `.coef_` attribute, we gain the coefficients of independent variables (x's) and with `.intercept_` attribute, we gain the intercept of our regression line

...

Ok; the regression line is created for our dependent variable (Age → y) and independent variables (other body measurements → x_0 to x_{11}).

Create regression line

To create regression line, at first, we should consider one of variables as dependent variable (y) and other variables as independent variables (x)

...

In this project, Age is y and other variables considered as x

...

`LinearRegression()` class, performs the regression operation for us;

With `fit()` method, we can fit our x's and y for regression operation as dependent and independent variables.

||| ----- Step 2: Make regression line ----- |||

The coefficient of variables are shown in bottom...

Gender: -2.1600681758606353

HeadCircumference: -0.29175212771841047

ShoulderWidth: 0.4913502191932724

ChestWidth: -0.2637862587951137

Belly: -0.043528822433128

Waist: 0.1583833238679603

Hips: 0.1655076168425505

ArmLength: 0.23630518352852511

ShoulderToWaist: 0.5897866941413992

WaistToKnee: 0.34539183792808575

LegLength: -0.06615631219156848

TotalHeight: 0.12851226045314823

Intercept of regression line is: -9.222430674200167

Average prediction error: 8.365045897383036 ~= 8

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (Predicted_i - Actual_i)^2}{N}}$$

Create regression line (results)

The results obtained from performing linear regression operations to predict age;

...

Coefficients of independent variables and intercept of regression line.

...

With predict() method, we can predict our desired option (which is Age here) based on the different independent variables; (shown in the previous slide)

...

To calculate error of our prediction, we use RMSE formula;

<< Root mean square error >>

← It's formula (predicted Ages – actual Ages)
N is the number of y's (from first age to the last).

Solve an example

In this part, we solve an example to check our prediction using the created equation (to predict age of a person with his/her body measurements);

...

In general we want to check obtained result of linear regression.



```
print("\n||| ----- Step 3: Solve an example ----- |||")
# You can use this numbers for example: 1, 20, 25, 24, 25, 36, 17, 23, 12, 9, 19, 48
x_for_predict = [[]]
for header in x_list:
    x_for_predict[0].append(int(input(f"Enter a number for input var <{header}>: ")))
predicted_example = reg_res.predict(x_for_predict)[0]
print("An example to prediction: ", abs(predicted_example), " ~= ", round(abs(predicted_example)), "years old")
```

```
||| ----- Step 3: Solve an example ----- |||
Enter a number for input var <Gender>: 1
Enter a number for input var <HeadCircumference>: 20
Enter a number for input var <ShoulderWidth>: 25
Enter a number for input var <ChestWidth>: 24
Enter a number for input var <Belly>: 25
Enter a number for input var <Waist>: 36
Enter a number for input var <Hips>: 17
Enter a number for input var <ArmLength>: 23
Enter a number for input var <ShoulderToWaist>: 12
Enter a number for input var <WaistToKnee>: 9
Enter a number for input var <LegLength>: 19
Enter a number for input var <TotalHeight>: 48
An example to prediction: 16.695157418056837 ~= 17 years old
```

Let's go to draw the plots

...


Here is the code for the charts that will be shown in the next slides;

Note: the regression line in the graphs, created for Age based on **one** variable; that is, **Age based on the per independent variables**; also actual age values and predicted values (that obtained from regression operation (with all independent variables)) will be show along with the regression line.

The Charts and Plots 😊

```
print("\n||| ----- Step 4: Draw the plots ----- |||")
for header in x_list:
    plt.grid()
    # New regression for per input
    reg_res = LinearRegression()
    reg_input = x_list[header].values.reshape(-1, 1)
    reg_res.fit(reg_input, y_list)
    y_reg_line = reg_res.coef_ * x_list[header] + reg_res.intercept_
    # The plot inputs
    plt.scatter(x_list[header], y_list, color="purple")
    plt.scatter(x_list[header], predicted_y, color="orange")
    plt.plot(x_list[header], y_reg_line, color="red")
    # The plot styles
    plt.xlabel(header, fontsize=15)
    plt.ylabel("Age", fontsize=15)
    plt.legend(["Original", "Predicted", "Regression line"])
    plt.show()
```



Note: All density ranges are considered for independent variables (not Age)



Age based on the Gender

Type of relation: indirect

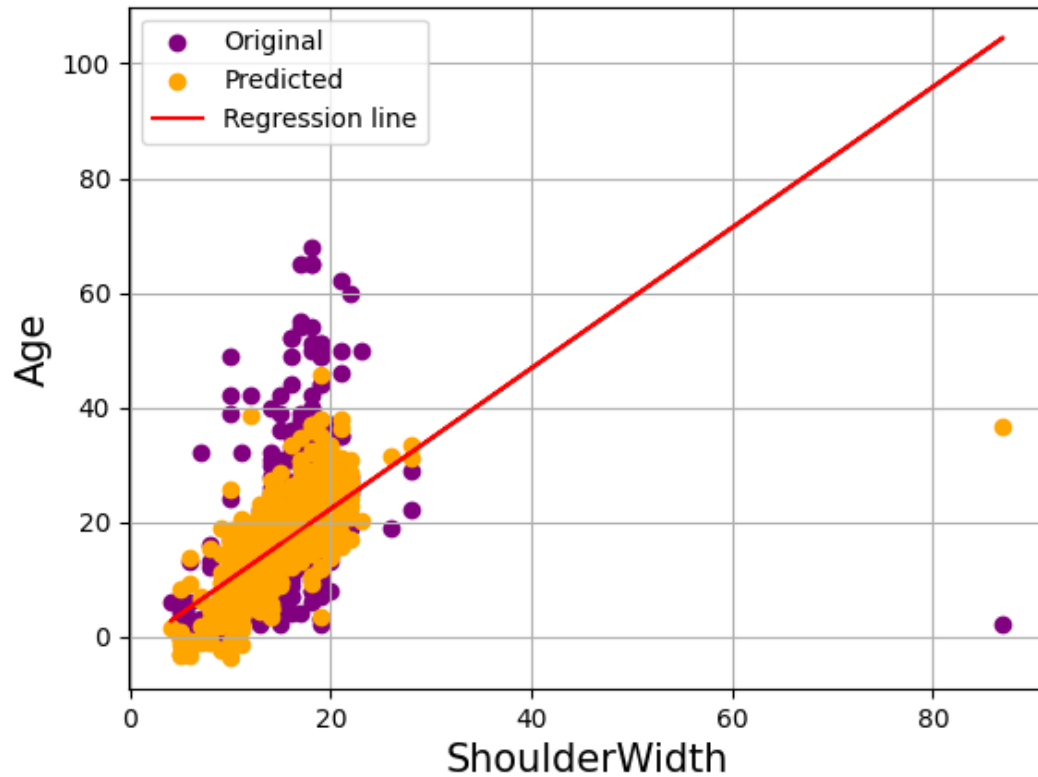
Gender 1 is male / Gender 2 is female



Age based on the Head circumference

Type of relation: direct

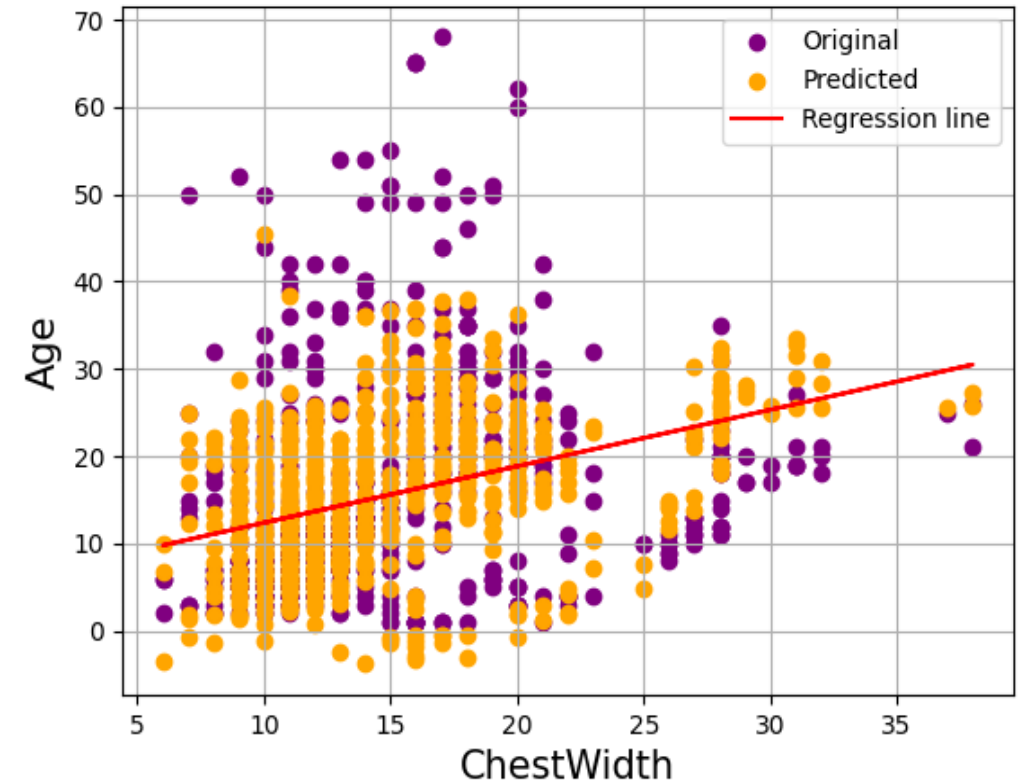
The highest density is between 10 and 30



Age based on the Shoulder width

Type of relation: direct

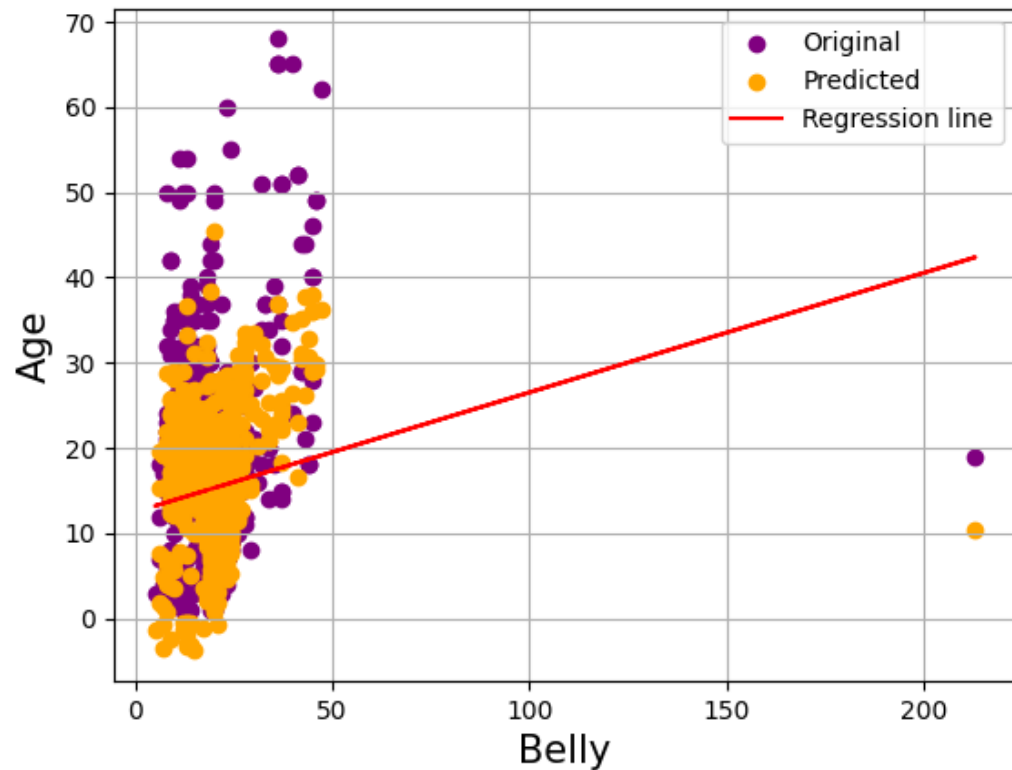
The highest density is almost between 4 and 25



Age based on the Head Chest width

Type of relation: direct

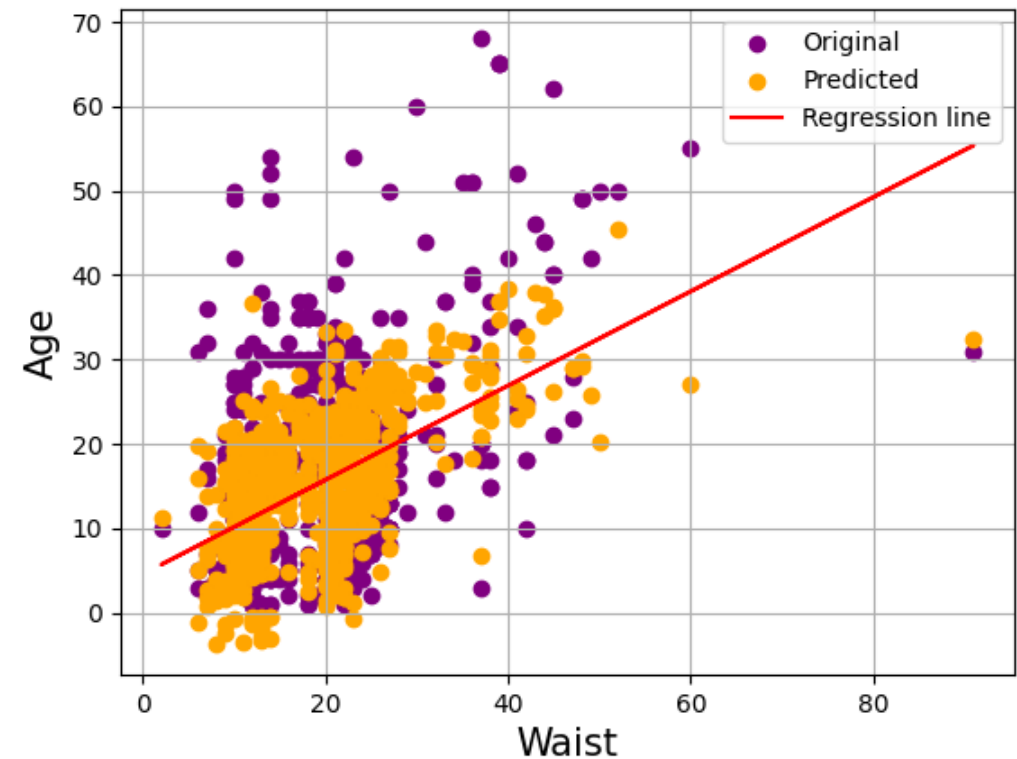
The highest density is between 5 and 25
(and a weaker density between 25 and 35)
In general, density is a little weak in this chart.



Age based on the Belly

Type of relation: direct

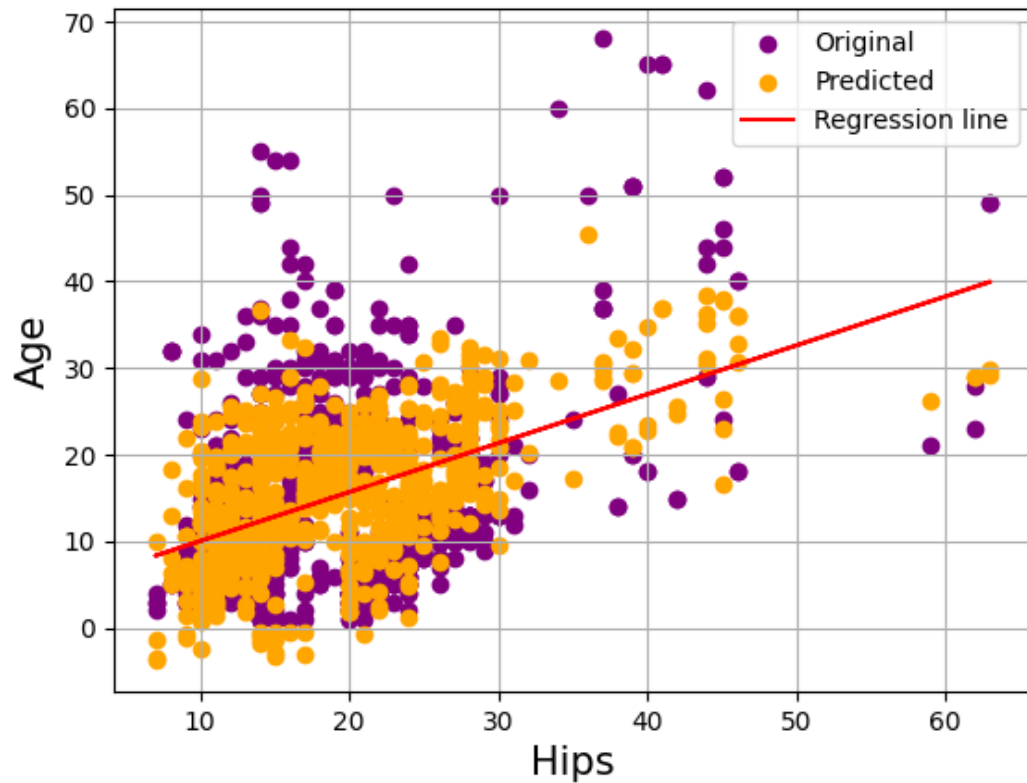
The highest density is between 0 and 50



Age based on the Waist

Type of relation: direct

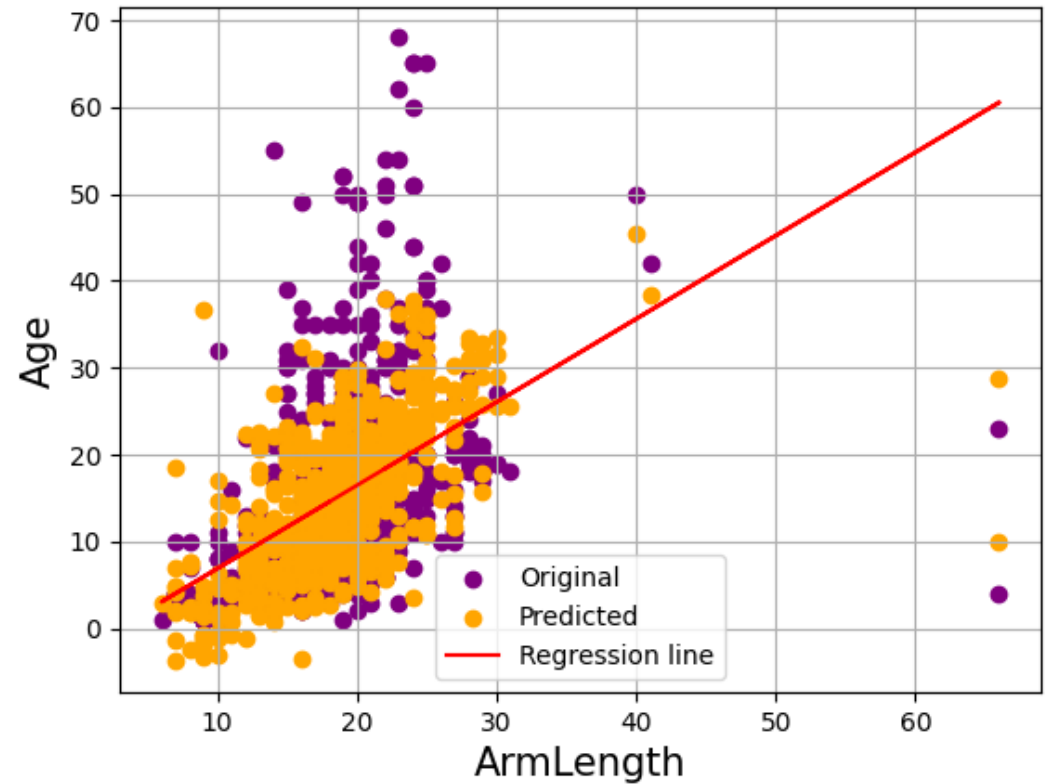
The highest density is between 0 and 25
(with a weaker density between 25 and 50 (to 60))



Age based on the Hips

Type of relation: direct

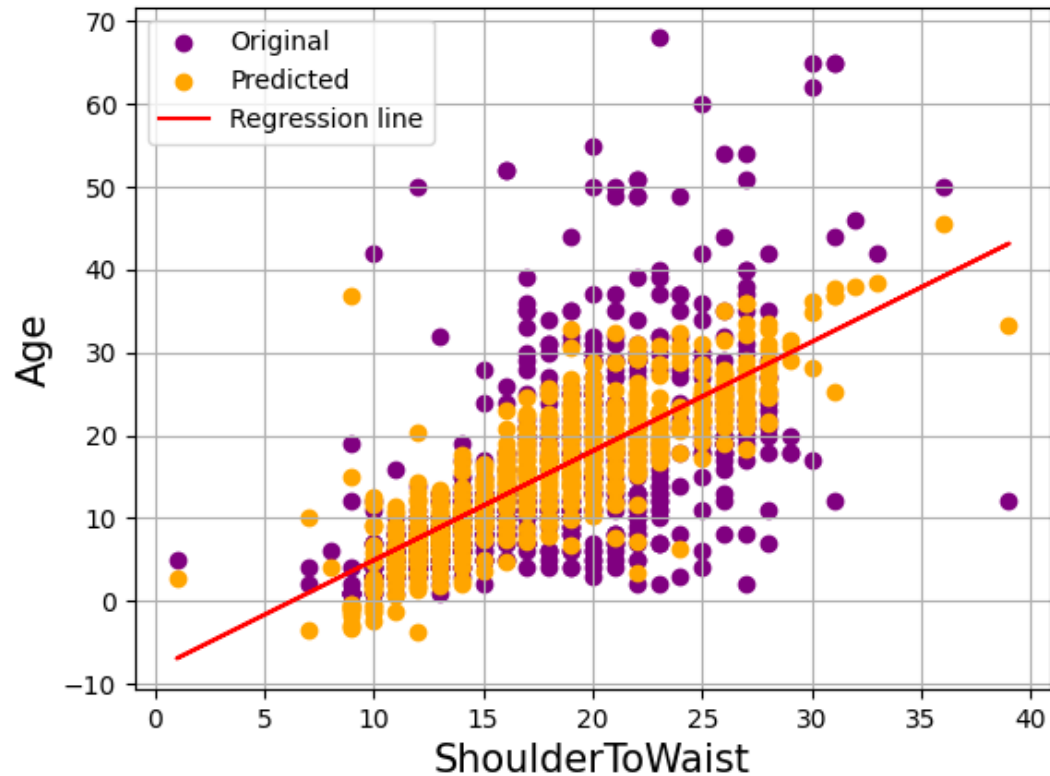
The highest density is between 4 and 30 (to 33)
(there is a weaker density between ~35 to ~57)



Age based on the Arm length

Type of relation: direct

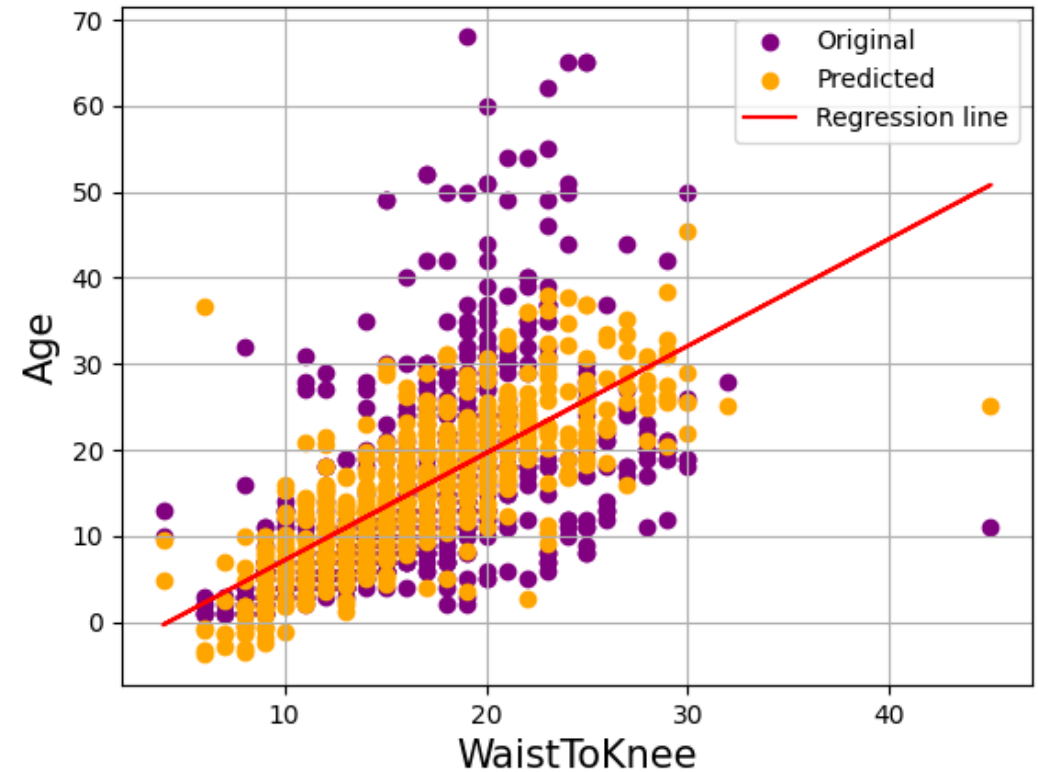
The highest density is between 4 and 30 (to 32)



Age based on the Shoulder to waist

Type of relation: direct

The highest density is between 5 (~6) and 30 (~28)

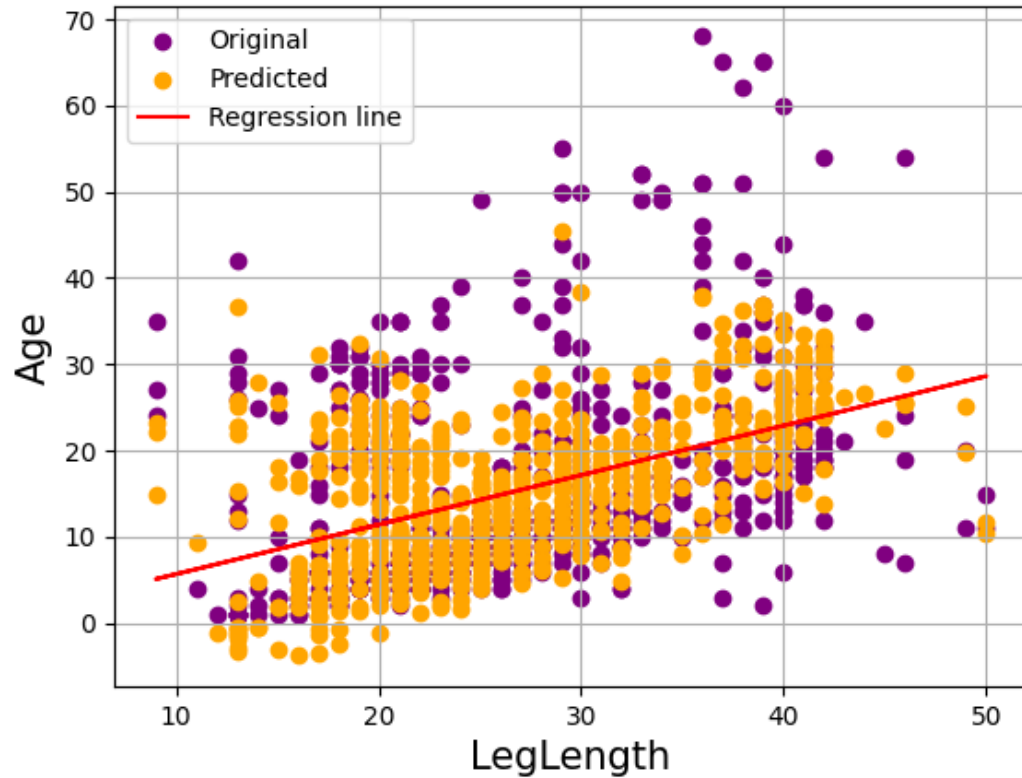


Age based on the Head Waist to knee

Type of relation: direct

The highest density is between 4 and 30 (~31)


The last slide



Age based on the Leg length

Type of relation: indirect

The highest density is between 10(~11) and 43
(with very weaker density to 50)



Age based on the Total height

Type of relation: direct

The highest density is almost between 15 and 75
(with very weaker density to 90)

The End

خدایا چنان کن سرانجام کار تو خشنود باشی و ما رستگار

By theMHD